Spatial detail falls (32 x 32 to 4 x 4) while width grows (32 to 256 channels).

Hidden tokens (hatched) are removed at stage 1 only, and F1 keeps their holes all the way to the decoder.

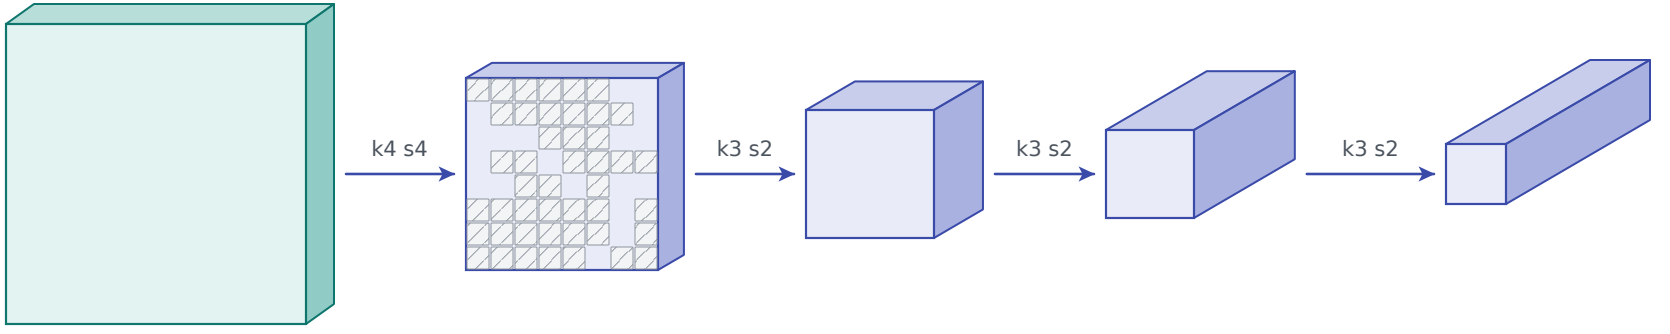

compressed input
(B, 32, 128, 128)

F1
(B, 32, 32, 32)
1/4 scale, 2 blocks, R=8

F2
(B, 64, 16, 16)
1/8 scale, 2 blocks, R=4

F3
(B, 160, 8, 8)
1/16 scale, 2 blocks, R=2

F4
(B, 256, 4, 4)
1/32 scale, 2 blocks, R=1

Front-face size follows the token grid; depth follows the channel count (both compressed for drawing). Each arrow is a patch embedding.