The model returns x_hat with exactly the input's shape; it never computes a loss or an anomaly score.

Then the traine

L1 + lambda(t) x SAM
on hidden, valid, eroded pixels

in raw reflectance space: an L1 of 0.008 means each band is off by 0.008
parts 23-25