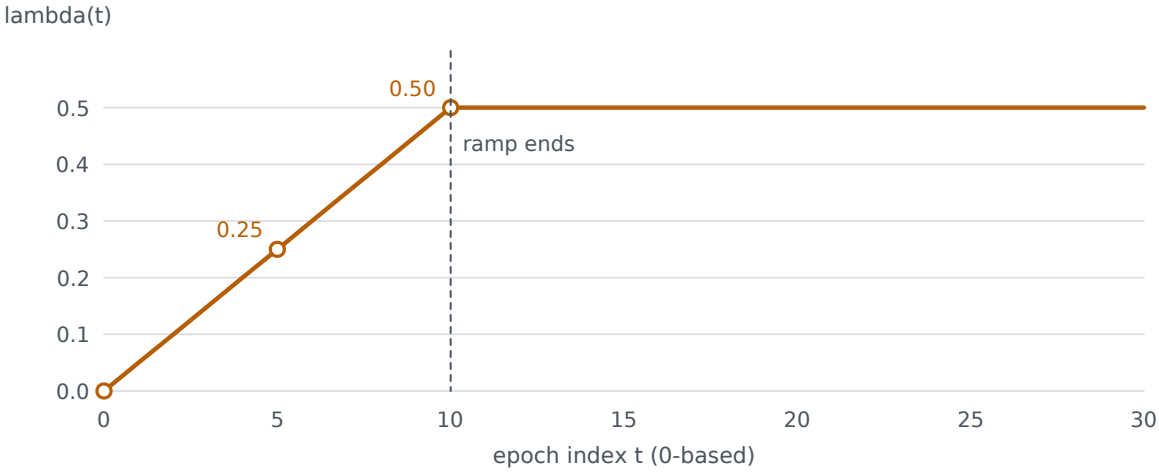

lambda(t) = 0.5 x min(1, t / 10)

loss = L1 + lambda(t) x SAM

Early on, SAM's gradients are noisy, so L1 leads; the angle term is phased in over 10 epochs.

v0.2.0: sam_weight 0.5, sam_ramp_epochs 10