One pixel, 202 bands

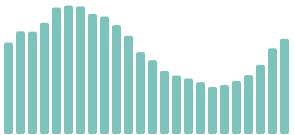

(202,)

band tokens
(202, 32)

value stretch + tag

K, V projections
(4, 202, 8) each

32 learned queries
(32, 32) -> Q

the same for every pixel

attention, 4 heads
32 x 202 scores/head

32 channels
o -> Linear(32,1)

Parameters, d = D = 32

value stretch	queries	Q, K, V, O	to_channel	**total**
64	1,024	4 x 1,056	33	**5,345**

No count depends on the band number: 202, 224 or 239 bands use the same 5,345 weights (today's conv: 5,376).