Trap 1 - hide pixels, not bands: a hidden pixel's tokens are removed before any spatial attention, so its reconstruction can only come from its neighbours.

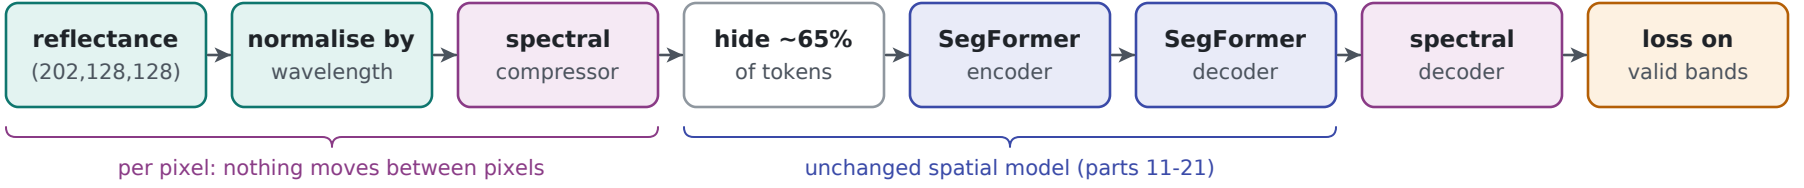

Trap 2 - normalise by a wavelength curve, never by per-scene min-max (it erases the absolute reflectance library matching needs).

Trap 3 - more sophisticated is not evidence: variants A, B, C are compared on reconstruction and on detection.