a Correcting the score map by edge strength

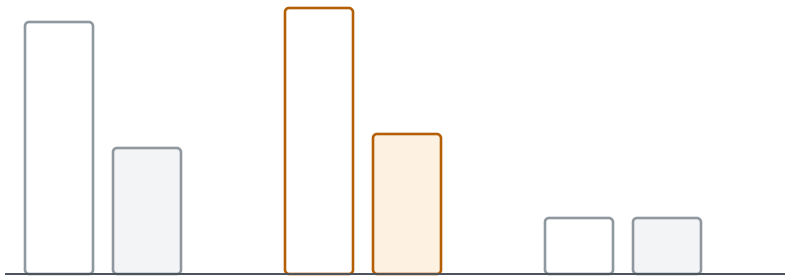

left bar: score before; right bar: after rank minus bin median

both strong-edge pixels drop together: bench 0.86 -> 0.67

49-68% of top-1% false positives on strong edges,
but also 29-79% of target pixels (chance: 10%)

(illustrative bar heights)

b Under the blind spot (one hidden 4 x 4 px token)

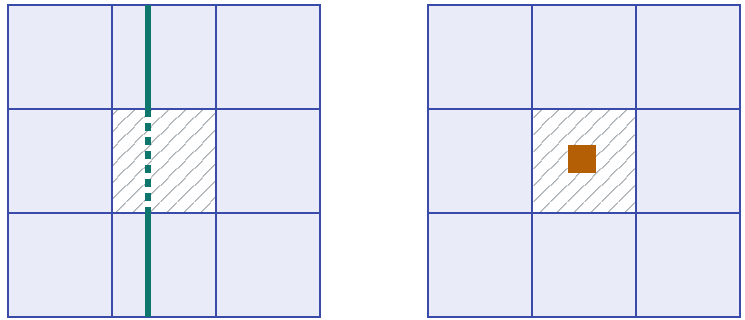

boundary: predictable **isolated target: not**

Every pixel is reconstructed while its own token is hidden
(two complementary checkerboard passes).

A gradient-matching loss on hidden pixels rewards
sharp continuation of what the context predicts.

Caveat: a target bigger than one token is partly
visible in the neighbouring tokens of the same pass.